



06 — Phase 2: Client Management System (Spec)

****Status:**** forward-looking design spec. Not built yet.
Target: ship behind a feature flag in Q3 2026; GA Q4 2026.





Contents

06 — Phase 2: Client Management System (Spec)

1. Goals

Anti-goals

2. Target users

3. The differentiated angle

3.1 Relationship Graph view

3.2 Hijri-aware account cadence

3.3 AI-suggested check-ins around Ramadan / Eid / National Days

3.4 Auto-generated QBR

3.5 Success Score model (transparent)

4. Key screens (described, not designed)

Screen A — Atlas Portfolio

Screen B — Account: Relationship Graph (default tab)

Screen C — Account: Health

Screen D — Account: QBR

Screen E — Cadences

Screen F — At-Risk Queue

5. Data model additions

subscriptions

usage_events

success_scores

stakeholders

stakeholder_relations

cadences

qbrs

commitments



Schema additions to existing tables

- 6. AI components
- 7. APIs to add
- 8. Integrations needed
- 9. Success metrics for Atlas itself
- 10. 90-day build outline
- 11. Risks & mitigations



Status: forward-looking design spec. Not built yet. Target: ship behind a feature flag in Q3 2026; GA Q4 2026.

Module name: NexFlow Atlas (working name) — the post-sale relationship intelligence layer.

Why this exists: every generic CRM can store a "customer" record after a deal closes. Almost none of them are post-sale operating systems and almost none of them treat a GCC client relationship as the multi-stakeholder, multi-generational, multi-channel, Hijri-calendar-aware engagement that it actually is. Atlas fills that gap.

■ 1. Goals

1. **One source of truth for "what is the relationship?"** Beyond contracts and tickets — relationship strength, decision-maker map, family / wasta context, last meaningful touch.
2. **Predict churn 90 days early.** A unified Success Score fed by support tickets, product usage, sentiment from logged calls, response latency, and contract milestones.
3. **Make every QBR write itself.** Auto-assemble a draft QBR doc from call intelligence, signals, and product usage.
4. **Be culturally fluent.** Hijri-aware account cadences. AI-suggested check-ins around Ramadan / Eid / National Days. Right-of-Center proper-name handling.
5. **Be Arabic-native.** Bilingual everywhere; client-facing artifacts (QBRs, emails, success summaries) generated in Arabic by default for KSA / GCC clients.

Anti-goals

- This module is not a help-desk replacement (we integrate with Zendesk / Freshdesk / Intercom rather than re-build).
- It is not a Customer Data Platform (segmentation lives in `/audiences`).



■ 2. Target users

ROLE	PRIMARY USE OF ATLAS
Customer Success Manager (CSM)	Daily home base. Health, action queue, QBR prep.
Account Executive	Pre-renewal expansion plays. Stakeholder map for upsell.
VP Sales / VP CS	Portfolio dashboard, churn risk forecast.
CEO	Top-20-account heatmap.
Implementation / Onboarding	Hand-off → Atlas timeline; sees go-live readiness (Phase 3).

■ 3. The differentiated angle

Five things competitors don't do well; Atlas does:

3.1 Relationship Graph view

A force-directed graph where nodes are **stakeholders** (not "contacts") and edges are **relationship signals**:

- "Reports to" / "Hired by" (org structure)
- "Champion" / "Detractor" / "Decision-maker" / "Influencer" / "Procurement gatekeeper"
- "Spoke to in the last 30 days" (auto-derived)
- "Family relation" (opt-in field, GCC context)
- "Same university / same regiment / same chamber" (wasta context, opt-in)

Each node carries a small avatar, a sentiment ring, and a "last touch" age halo. The graph is the **landing screen** of the account in Atlas.

3.2 Hijri-aware account cadence

Customer Success cadences (e.g. "QBR every 90 days", "monthly health check") respect the Hijri calendar:

- A scheduled monthly check-in shifts out of Ramadan working hours.
- A QBR is suggested in the week after Eid, not during it.
- National Day weeks are auto-blocked.
- "Last touch was Ramadan? Doesn't count toward dormancy timer."



3.3 AI-suggested check-ins around Ramadan / Eid / National Days

Atlas surfaces a **culturally-warm reach-out** prompt: - "It's the last Friday before Ramadan. Send the champion a short personal note." - "Eid Al-Fitr falls 14 days from now. Suggest a brief Eid greeting in Arabic with the customer's preferred dialect (KSA / Khaleeji)." - The AI drafts the message; CSM edits and sends in one click.

3.4 Auto-generated QBR

When a CSM opens "Prepare QBR" on an account, Atlas assembles in 30 seconds: - Executive summary (in client's preferred language) - Last 90 days: usage trend, NPS / CSAT pulse, support volume, key wins - Risk register, with mitigations - Roadmap of pending feature commitments - Suggested next-90-days plan - A printable, branded PPTX export (uses the same brand system as `docs/01-brand-palette-motion.md`)

QBR draft is a **structured document** the CSM edits in-place, not a static PDF.

3.5 Success Score model (transparent)

A **0-100** score per account, refreshed daily, with an **explanation card** so it's not a black box.

Inputs (weights tunable per org in Account Settings):

SIGNAL	DEFAULT WEIGHT	SOURCE
Product usage trend (DAU / WAU)	25%	<code>usage_events</code> (new)
Support ticket sentiment & volume	20%	help-desk integration
Call sentiment (last 90 days)	15%	<code>calls.sentiment_score</code>
Response latency (their side → ours)	10%	activity timeline
Champion engagement (last touch)	10%	activities + signals
Stakeholder coverage (≥ 3 stakeholders engaged)	10%	relationship graph
Renewal time horizon	5%	<code>subscriptions</code> (new)
Open invoices / payment health	5%	Quote-to-Cash

Each contribution is shown with a delta arrow ("+3 because product usage is up week-over-week"). When the score drops > 5 points in a week, Atlas creates a `notification` typed `ai`.



■ 4. Key screens (described, not designed)

All wireframe-level. Each respects the brand system (Living Mesh background, glass-panels, chameleon palette, motion tokens from docs/01).

Screen A — Atlas Portfolio

Default landing for CSMs and CS leaders. - Top KPI bar: NRR, GRR, churn risk count, expansion pipeline, NPS. - Heat-map grid of accounts colored by Success Score. - Side filters: segment, ARR band, owner, region, language. - Click an account → Screen B.

Screen B — Account: Relationship Graph (default tab)

- Big force-directed graph (zoom + pan).
- Right rail: account profile, current Success Score with explanation, ARR, renewal date, open commitments, outstanding tickets, next QBR date, **culturally-warm reach-out card** (§3.3).
- Tabs across the top: **Graph** | **Timeline** | **Health** | **Commitments** | **Tickets** | **QBR** | **Files**.

Screen C — Account: Health

- Success Score over time, with annotation pins for big events (CSM change, exec sponsor leaves, etc.).
- Per-signal breakdown chart.
- Recommended actions (AI-generated).

Screen D — Account: QBR

- Document editor with auto-generated sections.
- "Regenerate" button per section.
- "Export as PPTX / PDF" — branded.
- Comments / approval workflow with the AE.

Screen E — Cadences

- All recurring touches: monthly check-ins, quarterly QBRs, annual exec briefings.
- Each cadence shows next-fire date, Hijri-aware shifts, owner.
- Bulk-edit per segment.

Screen F — At-Risk Queue

- Triaged list of accounts with Success Score declining.
- Each row: top reason, suggested play, ETA to recovery.



- One-click "Open play" → starts the right workflow.

5. Data model additions

New Drizzle tables in `lib/db/src/schema/`. All `org_id`-scoped.

subscriptions

```
id, org_id, account_id (→ companies), plan, currency,  
arr_amount, started_at, renews_at, terms_url, autorenew,  
created_at, updated_at
```

usage_events

```
id, org_id, account_id, user_email, event_name,  
event_value (jsonb), occurred_at
```

(Ingested via API from the customer's product or via warehouse sync.)

success_scores

```
id, org_id, account_id, score (0–100), version,  
contributions (jsonb – per-signal deltas), captured_at
```

(One row per snapshot — daily.)

stakeholders

```
id, org_id, account_id, contact_id (→ contacts),  
role (champion | dm | influencer | gatekeeper | detractor),  
tenure_in_role_months, sentiment_label,  
notes, opt_in_relationship_context boolean,  
created_at, updated_at
```

stakeholder_relations

```
id, org_id, source_stakeholder_id, target_stakeholder_id,  
relation_type (reports_to | hired_by | family | university | chamber | other),  
opt_in boolean, notes,  
created_at
```


**cadences**

```
id, org_id, account_id, type (qbr | monthly | weekly | annual_exec),
next_fire_at, hijri_aware boolean, last_fired_at,
owner_id (→ users), active boolean
```

qbrs

```
id, org_id, account_id, status (draft | shared | signed_off),
document (jsonb – sections), generated_by (ai | human),
created_at, updated_at
```

commitments

```
id, org_id, account_id, title, owner_side (us | them),
due_at, status (open | done | overdue), source (call | email | qbr),
created_at, updated_at
```

Schema additions to existing tables

- **companies**: add **arr_amount** integer, **renewal_date** timestamp, **csm_id** uuid → users, **success_score** integer (denormalized latest), **success_score_updated_at**.
 - **notifications.notification_type**: extend the enum with **cs** (customer success), **qbr**, **cadence**.
-



6. AI components

COMPONENT	WHAT IT DOES	BACKING MODEL
Relationship-graph builder	Suggests new edges (champion / detractor / influencer) from call sentiment, email patterns, and meeting attendance.	Claude 3.5 Sonnet via OpenRouter.
Cultural cadence advisor	Decides when to shift / hold a touchpoint.	Rule-based (Hijri calendar) + LLM for nuance ("the customer's CFO is performing Hajj, push by 3 weeks").
Reach-out drafter	Drafts Eid / Ramadan / National Day greetings in customer's preferred language and dialect.	Claude 3.5 Sonnet, with a Saudi-Arabic style prompt.
Success Score explainer	Generates the human-readable explanation card.	gpt-4o-mini.
QBR auto-author	Generates each section from the underlying data (usage trend, sentiment, commitments).	Claude 3.5 Sonnet for narrative; deterministic compositor for structure.
At-risk play recommender	Picks the right play from the Sales Playbooks library + customizes to the account context.	gpt-4o-mini.

All call out via existing `aiChat` / `aiJson` in `artifacts/api-server/src/lib/ai.ts`. Atlas adds a new file `artifacts/api-server/src/lib/atlas.ts` for higher-level prompts.

7. APIs to add

Mounted under `/api/atlas`:

GET	<code>/accounts/:id/graph</code>	→ relationship graph nodes + edges
GET	<code>/accounts/:id/success-score</code>	→ latest score + explanation
GET	<code>/accounts/:id/success-score/history</code>	→ time-series
POST	<code>/accounts/:id/qbr/generate</code>	→ returns draft QBR document
PUT	<code>/accounts/:id/qbr/:qbrId</code>	→ save edits
POST	<code>/accounts/:id/qbr/:qbrId/share</code>	→ email it (Resend) and mark shared
GET	<code>/cadences</code>	→ list (filterable)
POST	<code>/cadences</code>	→ create
PUT	<code>/cadences/:id</code>	→ edit / shift
GET	<code>/atlas/portfolio</code>	→ CSM dashboard payload
GET	<code>/atlas/at-risk</code>	→ triaged at-risk queue



OpenAPI definitions added to `lib/api-spec/` ; codegen regenerates Zod schemas + React Query hooks for the web + mobile.

■ 8. Integrations needed

INTEGRATION	WHAT WE USE IT FOR	NOTES
Help-desk (Zendesk / Freshdesk / Intercom)	Ticket sentiment + volume into Success Score.	Use connector pattern from <code>query-integration-data</code> skill.
Product usage (warehouse: BigQuery / Snowflake / Databricks; or direct webhook)	DAU / WAU / event counts.	Ship a NexFlow JS SDK for direct event ingest as well.
Hijri calendar	Cadence shifting.	Use <code>umm-al-qura</code> calendar lib (npm).
Translation / Arabic NLP	Language polish on AI drafts.	Goes through OpenRouter (Claude 3.5 handles Arabic well).
Payment health	Open invoices into Success Score.	Reuse Quote-to-Cash data.

■ 9. Success metrics for Atlas itself

METRIC	TARGET (12 MO. POST-GA)
% of paying NexFlow customers with Atlas turned on	≥ 70%
Avg Success Score lift across active orgs	+12 points y/y
Churn rate among Atlas-tracked accounts	≤ 6% (vs. 10% control)
QBRs auto-generated per CSM per quarter	≥ 8
Time saved per QBR	≥ 4 hours
% of QBRs sent in Arabic to GCC clients	≥ 60%
NPS of CSMs using Atlas	≥ 50



10. 90-day build outline

Two-week sprints, two engineers + one designer + 0.5 PM.

SPRINT	SCOPE
S1 (wks 1-2)	Spec lock, data model PR, migration, seed sample data. Read-only <code>companies.arr_amount</code> + <code>csm_id</code> fields shipped.
S2 (wks 3-4)	Success Score v1 (signals: call sentiment, response latency, ticket count via help-desk webhook). Daily cron snapshots. Score on company detail page.
S3 (wks 5-6)	Cadences module + Hijri awareness. Monthly check-in cadence ships. Atlas Portfolio screen v1.
S4 (wks 7-8)	Relationship Graph (read-only). Stakeholder roles + manual edit. Account detail tabs scaffold.
S5 (wks 9-10)	QBR auto-author v1 (structured doc). PPTX export reusing brand system. Reach-out drafter for Ramadan / Eid.
S6 (wks 11-12)	At-Risk Queue + recommended plays. Mobile app: Atlas tab with Portfolio + account detail. Internal beta with two design-partner customers.

GA after a 6-week design-partner phase and a single dedicated hardening sprint.

11. Risks & mitigations

RISK	MITIGATION
Customers don't use opt-in family / wasta fields	Make them optional and clearly value-add. Lead with org-structure edges.
Hijri calendar bugs around lunar visibility	Defer hard cutoffs to Umm al-Qura; allow per-account override.
LLM-drafted QBRs feel generic	Train on top-quartile real QBRs; CSM always edits before send.
Success Score gamed	Show the contribution breakdown; orgs can re-weight.
Help-desk integration sluggish	Cache last 24 h of ticket signals; fall back to manual entry.